

1) Executive Summary

What we're building: Earnings Call Insight (ECI) is a web app that helps users quickly understand what management cared about over time by analyzing quarterly earnings call Q&A (investor questions + management answers) and annual reports. Users choose a ticker and a time window (year + quarter); ECI retrieves the relevant files from our internal document store, routes them through multiple LLMs (e.g., GPT-5, Claude by Anthropic, Gemini by Google), and returns an insight summary: top themes, turning points, and cited evidence from the source docs.

Why now? Analysts and operators drown in text: multiple years \times 4 quarters/year \times long transcripts. Existing tools summarize single calls, but rarely track how focus shifts (e.g., from “growth” to “margin/FCF”) with clear citations. ECI solves this by (1) structuring content around Q&A; (2) extracting themes across quarters; and (3) returning explainable, evidence-linked outputs.

What success looks like: A user can select *AMZN, 2022Q1–2023Q4* and, within seconds, see Top 5 themes with quarter-level heat, cited quotes, and detected inflection points; download/share results; and compare to peers.

2) User Research & Problem Definition

Personas (2–3)

1. **Sell-side equity analyst**
 - Needs: fast quarterly deltas; defensible theses; source citations for notes.
 - Pain: hours scanning transcripts; inconsistent tagging; misses subtle shifts.
2. **Buy-side analyst**
 - Needs: high-signal summaries before investment decisions; YoY/QoQ focus changes; risk red flags.
 - Pain: fragmented tools; manual spreadsheets and analysis; hard to compare peers.
3. **Academic researcher**
 - Needs: learn a sector’s history fast; exportable graphs; glossary-level clarity.
 - Pain: steep ramp-up; “too much text, too little structure”.

Current vs. Ideal Journeys

- Current way to do it: Search transcripts in database (capital IQ, company website, etc) \rightarrow skim 60 - 90 min texts \rightarrow CTRL+F search for keywords \rightarrow paste into notes \rightarrow subjective highlights \rightarrow weak quarter-over-quarter continuity.
- Ideal (ECI solution): Pick ticker + time range \rightarrow ECI retrieves relevant text from database \rightarrow multi-LLM analysis \rightarrow theme timeline + top quotes + changes over time with citations \rightarrow export/share.

Problem Validation (evidence)

Validated by our team members' personal experience and lots of inconvenience we observed when reading dozens of calls per company per year. When analysts are doing equity research, most of the time is spent locating and skimming, and most tools summarize one call; It always needs trend + evidence across quarters summarized by themselves manually, because existing information platforms like Capital IQ and FactSet, don't deliver quarterly trend tracking with source-linked proof.

3) Product Vision (MVP scope & success metrics)

MVP scope (Sprint 2–3):

- Upload/store PDF/TXT earnings calls & annuals per ticker/quarter.
- Quarter-aware selection UI: {ticker, start:{y,q}, end:{y,q}}.
- Backend service (Django) that accepts doc path, reads text, runs multi-model LLM routing (GPT-5 / Claude / Gemini), and returns:
 - Top themes (3–5) with quarter heat;
 - Turning points analysis if any (e.g., “2023Q3: focus shifts to FCF”).
 - Risks/uncertainties with evidence and citations.

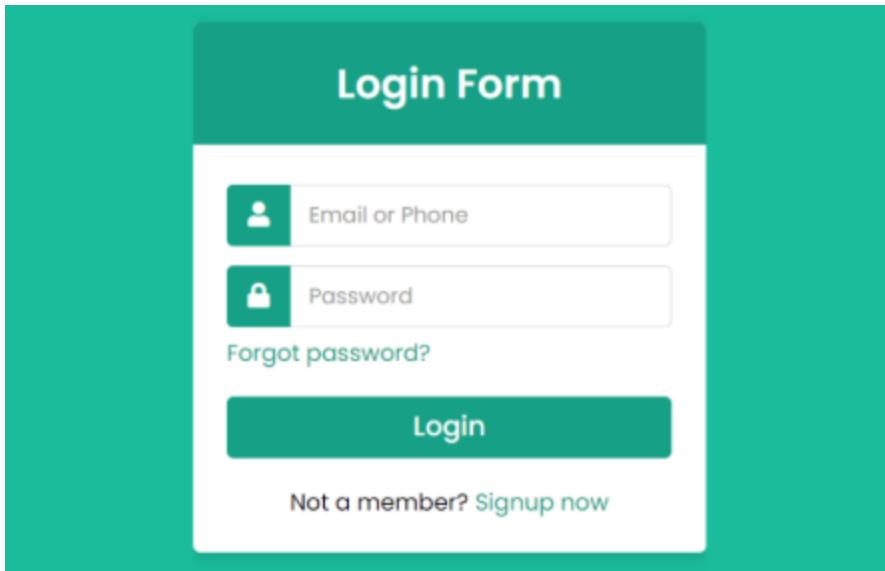
Nice-to-have (post-MVP): peer comparison; CSV/PPT export; topic alerts; quantitative KPI extraction.

Success metrics (first semester):

- Accuracy and user trust: users can have a smooth experience retrieving the information needed through ECI. The information can be useful and complete enough for their real-world analysis at work.
- Coverage: ≥ 2 tickers \times 4 years (16 quarters) each.
- Information collection: through registration page, collect users information (name, email address) successfully
- Adoption: ≥ 10 active users (class/club partners).

4) Wireframes / Mockups

The login page will look similar to the pic below.



The page for ticker/time selection and output will look similar to the pic below.

Quarterly Data Selection

Company Name Starting Quarter Ending Quarter

Select Company... Q1 2023 Q4 2023 Submit

Output Data

5) Technical Architecture (initial decisions)

Backend:

Django (chosen for stability, batteries-included admin, easy REST endpoints).

LLM Orchestration:

LiteLLM (one API for GPT-5/4o depending on cost/time balance, Claude, Gemini).

Prompting:

Task-specific templates (Trend / Investor-Focus / Claim-Check), JSON-only outputs, evidence IDs required.

Data store:

File stored in the system folder directly given that we currently have a comparably small size of files. We can set up optional SQLite for metadata (MVP) later if we have capability.

Retrieval:

MVP concatenates targeted files; next sprint migrates to RAG (segmenting Q&A, embeddings, ANN search).

Deployment:

Dev on local/Colab; stage on Render.

6) Product Backlog Summary

- **Stories: 10**
- **Total story points: 42** (Fibonacci)
- **Priorities: High / Medium / Low**

Examples:

Stories	Acceptance Criteria / AC	Story Points / SP	Priority / Label
<i>As a user, I can select ticker + quarter range so that results focus on the exact period.</i>	valid ticker; invalid quarter rejected; range inclusive; persisted in URL.	3	High
<i>As a router service, I can pass doc_paths[] to the analysis API so the backend reads only those files.</i>	rejects paths outside allowed folder; returns 400 on missing file; logs used paths.	5	High
<i>As a user, I see Top 5 themes with quarter-by-quarter heat and one or more citations per theme.</i>	JSON schema validated; ≥1 quote/file path per theme; handles no-evidence gracefully.	8	High

<i>As a PM/analyst, I see turning points detected with linked quotes.</i>		5	High
<i>As a developer, I can switch LLM models per request.</i>		5	Medium
<i>As a user, I can open an evidence drawer to view quotes with file path and quarter.</i>		5	Medium
<i>As a system, I reject inputs without ticker/start/end/doc_paths with clear 400 errors.</i>		2	High
<i>As a system, I limit token usage by truncating overly long files and warn in the payload.</i>		3	Medium
<i>As a developer, I have a JSON-only LLM schema for themes/turns/risks.</i>		3	High
<i>As a user, I can export the JSON results.</i>		3	Low
<i>Total Story Points</i>			

7) Sprint 2 Plans

Sprint goal: Get a working web app that can take a ticker and quarter range, run analysis, and return basic insights.

User stories selected in sprint 2: No.1- No.5 stories in Part 6.

Total story points: 26 pts

Team member assignment:

Alex Xu → front-end

Peter Yang → backend

Sarah Zhang → integration and testing

Time committed per week per member: 3-4hrs each + 3hrs together

Identified risks:

LLM API limits access; Long files may cause slow responses.

8) Team Organization

Link to team charter, roles, meeting schedule:

- Project name: Randomwalkwallstreet
- Team members: Alex Xu, Sarah Zhang, Peter Yang
- Roles: Scrum Master: rotates each sprint; Product Owner Proxy: Alex Xu; Developers: Sarah Zhang, Peter Yang.
- Communication: Email, Wechat and zoom calls; Weekly meeting

- Working Agreements: Code reviews within three days; Branch naming to be determined; Commit format short and clear
- Role Rotation: Scrum Master changes every sprint (Alex → Sarah → Peter → Alex...)
- Conflict Resolution: Talk directly and respectfully; Discuss as a team if needed; Scrum Master helps mediate.

9) Process Setup

Link to GitHub Project:

github.com/Atmoskun/randomwalkwallstreet

Link to Definition of Done:

github.com/Atmoskun/randomwalkwallstreet/blob/main/Randomwalk/docs/definition-of-done.md

Link to Definition of Ready:

github.com/Atmoskun/randomwalkwallstreet/blob/main/Randomwalk/docs/definition-of-ready.md